

→ Gradient Descent


- Have some function $J(w, b)$

- Want $\min_{w, b} J(w, b)$

- Outline:

 - Start with some w, b

 - Keep changing w, b to reduce $J(w, b)$ until we settle at or near a minimum (may have > 1 min)



→ Implement Gradient Descent Algorithm

$$w = w - \alpha \frac{d}{dw} J(w, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b)$$

} Update steps

↳ Simultaneously update w and b

↳ Repeat until converge

$\alpha \equiv$ Learning Rate

↳ Controls the step size

- Learning Rate

→ If α is too small \Rightarrow Gradient descent may be slow

→ If α is too large \Rightarrow Gradient descent may:

- overshoot, never reach minimum

- fail to converge

- Near a local minimum

→ Derivatives become smaller

→ Update steps become smaller

} Can reach minimum without decreasing learning rate α

→ Gradient descent for linear regression

- Linear Regression model

$$f_{w, b}(x) = wx + b$$

- Cost Function

$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m [f_{w, b}(x^{(i)}) - y^{(i)}]^2$$

- Gradient Descent Algorithm

$$w = w - \alpha \frac{\partial}{\partial w} J(w, b) \quad \Leftrightarrow \quad w = w - \alpha \frac{1}{m} \sum_{i=1}^m [f_{w, b}(x^{(i)}) - y^{(i)}] x^{(i)}$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w, b) \quad \Leftrightarrow \quad b = b - \alpha \frac{1}{m} \sum_{i=1}^m [f_{w, b}(x^{(i)}) - y^{(i)}]$$